

p8106_stl2137_hw4

Question 1

```
### Loading in prostate data
data("Prostate")
dat_prostate <- Prostate
```

Part A

```
set.seed(1)

tree_prost_1 <- rpart(formula = lpsa ~., data = dat_prostate)
rpart.plot(tree_prost_1)

### Finding lowest cp
cp_table <- printcp(tree_prost_1)

##
## Regression tree:
## rpart(formula = lpsa ~ ., data = dat_prostate)
##
## Variables actually used in tree construction:
## [1] lcavol  lweight  pgg45
##
## Root node error: 127.92/97 = 1.3187
##
## n= 97
##
##          CP nsplit rel error  xerror    xstd
## 1 0.347108      0   1.00000 1.01323 0.162162
## 2 0.184647      1   0.65289 0.88779 0.111915
## 3 0.059316      2   0.46824 0.59168 0.066102
## 4 0.034756      3   0.40893 0.61359 0.069269
## 5 0.034609      4   0.37417 0.58640 0.067630
## 6 0.021564      5   0.33956 0.57853 0.068772
## 7 0.021470      6   0.31800 0.56398 0.067155
## 8 0.010000      7   0.29653 0.54721 0.068034


#plotcp(tree_prost_1)

minErr <- which.min(cp_table[,4])

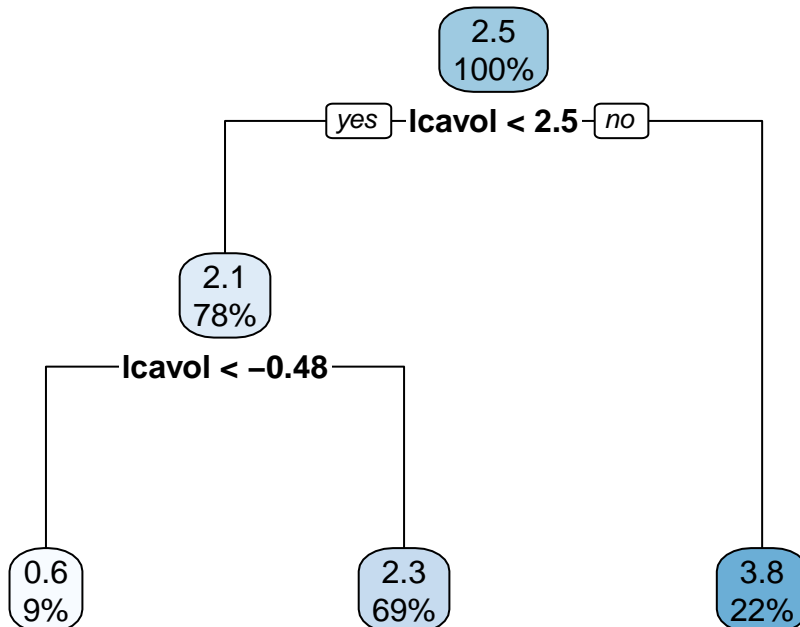
# minimum cross-validation error
tree_prost_3 <- prune(tree_prost_1, cp = cp_table[minErr, 1])

# 1SE rule
tree_prost_4 <- prune(tree_prost_1, cp =
  cp_table[cp_table[,4] < cp_table[minErr, 4] + cp_table[minErr, 5], 1][1])
```

```
rpart.plot(tree_prost_3)
```




```
rpart.plot(tree_prost_4)
```



The tree size that corresponds to the lowest cross-validation error is 8. This is not the same tree size as the one obtained using the 1 SE rule, as the tree size is 3.

Part B

```
plotcp(tree_prost_1)
```




Based off the leftmost value for which the mean lies below the horizontal line in the cp plot, a $cp = 0.1$ and a tree size of 3 should be utilized.

```
set.seed(1)

final_tree_prost <- rpart(lpsa ~ ., data = dat_prostate,
                          control = rpart.control(cp = 0.1))

rpart.plot(final_tree_prost)
```



If you have a log cancer volume greater than 2.5, we predict that you will have a log prostate specific antigen level of 3.8.

Part C

```
set.seed(1)

bagging_prost <- randomForest(lpsa ~ ., data = dat_prostate,
                              mtry = 8)
```

```
bagging_prost$importance
```

```
##          IncNodePurity
## lcavol      76.557359
## lweight    16.761566
## age        5.875410
## lbph       5.123664
## svi        6.534788
## lcp        5.937665
## gleason    1.096503
## pgg45      5.776304
```

The variable for log cancer volume, at a value of 76.557359, has the highest variable importance. The variable for log prostate weight, at the value of 16.7615664, has the 2nd highest variable importance.

Using caret for RF


```
control <- trainControl(method = "cv")

rf_grid_prost <- expand.grid(mtry = 1:8,
                             splitrule = "variance",
                             min.node.size = 1:8)
```

```
set.seed(1)
rf_fit_prost <- train(lpsa ~ ., dat_prostate,
  method = "ranger",
  tuneGrid = rf_grid_prost,
  trControl = control)
```

```
ggplot(rf_fit_prost, highlight = TRUE)
```

```
## Warning: The shape palette can deal with a maximum of 6 discrete values
## because more than 6 becomes difficult to discriminate; you have 8.
## Consider specifying shapes manually if you must have them.
## Warning: Removed 16 rows containing missing values (geom_point).
```



Part D

```
set.seed(1)
rf_prost <- randomForest(lpsa ~ ., data = dat_prostate,
  mtry = 3)
```

```
rf_prost$importance
```

```
##          IncNodePurity
## lcavol      44.037365
## lweight     20.439936
## age         7.901313
```

```
## lbph      6.974511
## svi       12.779901
## lcp       13.149594
## gleason   4.205354
## pgg45     11.130664
```


The variable for log cancer volume, at a value of 44.037365, has the highest variable importance. The variable for log of prostate weight, at the value of 20.439936, has the 2nd highest variable importance.

Caret for Boosting

```
gbm_grid <- expand.grid(n.trees = c(2000,3000),
                       interaction.depth = 2:10,
                       shrinkage = c(0.001,0.003,0.005),
                       n.minobsinnode = 1)

set.seed(1)
gbm_fit_prost <- train(lpsa ~ ., dat_prostate,
                      method = "gbm",
                      tuneGrid = gbm_grid,
                      trControl = control,
                      verbose = FALSE)


ggplot(gbm_fit_prost, highlight = TRUE)
```




Part E

```
set.seed(1)
boosting_prost <- gbm(lpsa ~ ., data = dat_prostate,
  distribution = "gaussian",
  n.trees = 5000,
  interaction.depth = 3,
  shrinkage = 0.005,
  cv.folds = 10)

nt <- gbm.perf(boosting_prost, method = "cv")
```



```
boosting_plot <- summary(boosting_prost, las = 2, cBars = 19, cex.names = 0.6)
```



The variable for log cancer volume, at a value of 34.229525, has the highest variable importance. The variable for log of prostate weight, at the value of 23.1311753, has the 2nd highest variable importance.

Part F

```
resamp <- resamples(list(rf = rf_fit_prost, gbm = gbm_fit_prost))
summary(resamp)
```

```
##
## Call:
## summary.resamples(object = resamp)
##
## Models: rf, gbm
## Number of resamples: 10
##
## MAE
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## rf  0.4737468 0.5269883 0.6122238 0.6001549 0.6672875 0.7080113    0
## gbm 0.4988605 0.5382599 0.6238535 0.6060614 0.6541746 0.7189084    0
##
## RMSE
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## rf  0.5831535 0.6354143 0.7101871 0.7381295 0.8276041 0.9995138    0
## gbm 0.6127814 0.6613586 0.7369285 0.7499428 0.8330904 0.9285269    0
##
## Rsquared
##      Min.   1st Qu.   Median     Mean   3rd Qu.     Max. NA's
## rf  0.3927114 0.5134677 0.5953671 0.6059044 0.7308171 0.7825008    0
## gbm 0.3912307 0.4940316 0.6253058 0.6103438 0.7211938 0.8171347    0
```

I would pick the random forest model. When comparing the RMSE of the two models when resampling, random forest has the lower mean RMSE.

Problem 2

Pulling & Creating Training/Test Datasets

```
data("OJ")
dat_oj <- OJ %>%
  janitor::clean_names()

set.seed(1)

rowTrain <- createDataPartition(y = dat_oj$purchase,
                                p = 799/1070,
                                list = FALSE)

train_dat_oj <- as.data.frame(dat_oj[rowTrain,])
test_dat_oj <- as.data.frame(dat_oj[-rowTrain,])
```


Part A

```
set.seed(1)
oj_tree <- rpart(purchase ~., data = train_dat_oj,
                 control = rpart.control(cp = 0))
#rpart.plot(oj_tree)

cp_table_oj <- printcp(oj_tree)

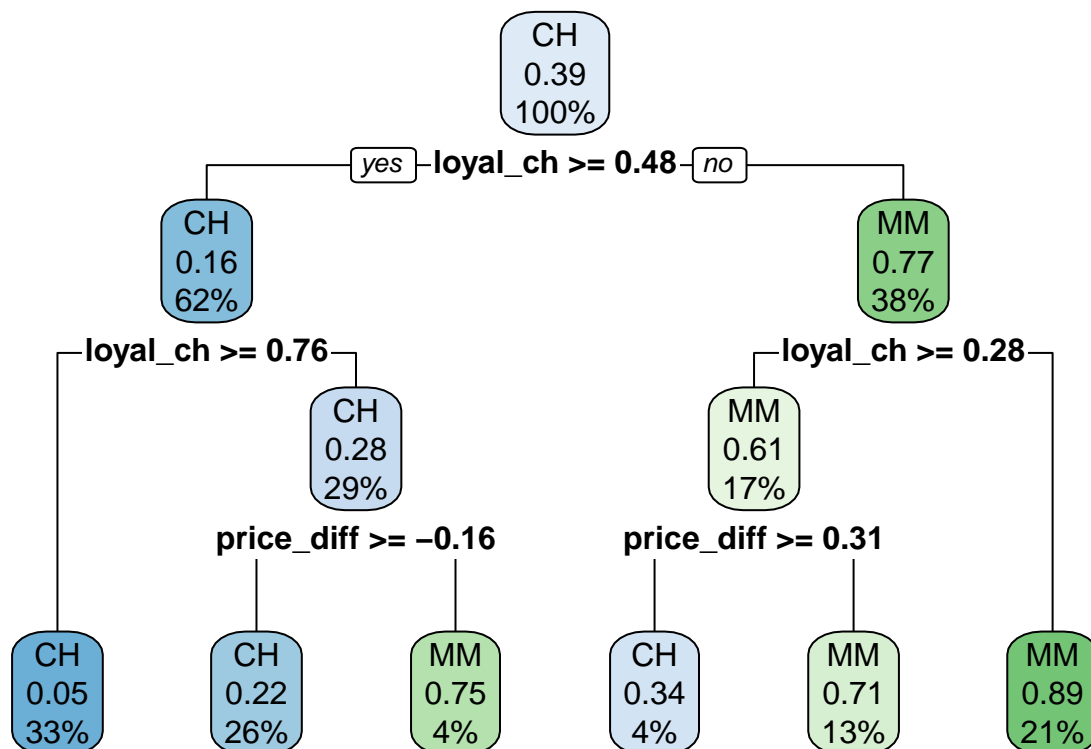
##
## Classification tree:
## rpart(formula = purchase ~ ., data = train_dat_oj, control = rpart.control(cp = 0))
##
## Variables actually used in tree construction:
## [1] list_price_diff loyal_ch      price_ch      price_diff
## [5] store          weekof_purchase
##
## Root node error: 312/800 = 0.39
##
## n= 800
##
##      CP nsplit rel error  xerror    xstd
## 1  0.51923077     0  1.00000 1.00000 0.044217
## 2  0.02243590     1  0.48077 0.48718 0.035564
## 3  0.01762821     3  0.43590 0.48397 0.035474
## 4  0.00801282     5  0.40064 0.48077 0.035384
## 5  0.00534188     7  0.38462 0.53526 0.036843
## 6  0.00480769    10  0.36859 0.56090 0.037477
## 7  0.00320513    14  0.34936 0.55449 0.037321
## 8  0.00160256    16  0.34295 0.57051 0.037706
## 9  0.00080128    18  0.33974 0.57372 0.037781
## 10 0.00000000    22  0.33654 0.57051 0.037706

plotcp(oj_tree)
```



```
minErr_oj <- which.min(cp_table_oj[,4])

# minimum cross-validation error
oj_tree_2 <- prune(oj_tree, cp = cp_table_oj[minErr_oj, 1])
rpart.plot(oj_tree_2)
```



Building Confusion Matrix

```

test_oj_tree_prob <- predict(oj_tree_2, newdata = test_dat_oj, type = "prob")
test_oj_tree_prob <- test_oj_tree_prob[,1]
test_oj_tree_pred <- rep("CH", 270)
test_oj_tree_pred[test_oj_tree_prob > 0.5] <- "MM"

oj_tree_matrix <- caret::confusionMatrix(data = as.factor(test_oj_tree_pred),
  reference = test_dat_oj$purchase,
  positive = "MM")

```

Based off the confusion matrix, it has an accuracy of 0.1740741. Thus, it has an error of 0.8259259.


Part B

```

set.seed(1)
rf_oj <- ranger(purchase ~., train_dat_oj,
  mtry = 6,
  min.node.size = 5,
  splitrule = "gini",
  importance = "permutation",
  scale.permutation.importance = TRUE)

barplot(sort(ranger::importance(rf_oj), decreasing = FALSE),
  las = 2, horiz = TRUE, cex.names = 0.7,
  col = colorRampPalette(colors = c("cyan", "blue"))(8))

```




The customer brand loyalty for Citrus Hill orange juice, at a value of 108.8103408, has the highest variable importance. As you can see in the bar plot, it is much higher variable importance compared to the other variables.

Part C


```
set.seed(1)

boosting_oj <- gbm(purchase ~., data = train_dat_oj,
  distribution = "gaussian",
  n.trees = 5000,
  interaction.depth = 3,
  shrinkage = 0.005,
  cv.folds = 10)

nt_oj = gbm.perf(boosting_oj, method = "cv")
```




```
summary(boosting_oj, las = 2, cBars = 19, cex.names = 0.6)
```



```
##           var    rel.inf
## loyal_ch      loyal_ch 56.8781742
## weekof_purchase weekof_purchase 9.9104369
## price_diff      price_diff 9.8463273
## list_price_diff list_price_diff 4.4288931
## store_id        store_id 3.4056419
## sale_price_mm    sale_price_mm 2.9345853
## store           store 2.1489733
## price_mm         price_mm 1.9319370
## price_ch         price_ch 1.7584702
```

```
## sale_price_ch      sale_price_ch  1.5642614
## disc_mm            disc_mm      1.3409564
## special_ch         special_ch   0.9929761
## disc_ch            disc_ch      0.8858828
## pct_disc_mm        pct_disc_mm  0.8351668
## special_mm         special_mm   0.7550823
## pct_disc_ch        pct_disc_ch  0.2361005
## store7             store7       0.1461343
```

```
boosting_imp_oj <- summary(boosting_oj)
```



The customer brand loyalty for Citrus Hill orange juice, at a value of 56.8781742, has the highest variable importance. As you can see in the bar plot, it is much higher variable importance compared to the other variables. The week of purchase variable, at a value of 9.9104369, has the 2nd highest variable importance.